

# Securing Kubernetes with Open Policy Agent (OPA) and Gatekeeper

Anton Sankov, 03.10.2022

# Anton Sankov

- Senior Software Engineer @ VMware Carbon Black
- Passionate about Kubernetes Security



a\_sankov



Anton Sankov



<https://asankov.dev>



# Agenda

Why Kubernetes Security

Different aspects of Kubernetes Security

Implementing Security

Open Policy Agent (OPA) and Gatekeeper

Demo

# Why Kubernetes Security?

**Kubernetes is more than orchestrator,  
It is a platform**

# Kubernetes as a Platform

- Orchestrate containers and ensure availability
  - Orchestrate storage and other resources (Ingress, certificates, etc.)
  - Secrets and configuration management
- 
- Kubernetes for CI/CD pipelines
  - Kubernetes for managing Kubernetes

**A system is only as secure as its  
weakest link**

# Kubernetes Exploits

An attacker that got access to your cluster can:

- Deploy malicious workloads/crypto miners
- Steal/Tamper with your data/secrets
- Do a container escape and get access to your whole infrastructure

9:45 AM

**Hacking Kubernetes: Live Demo Marathon**

Andrew Martin

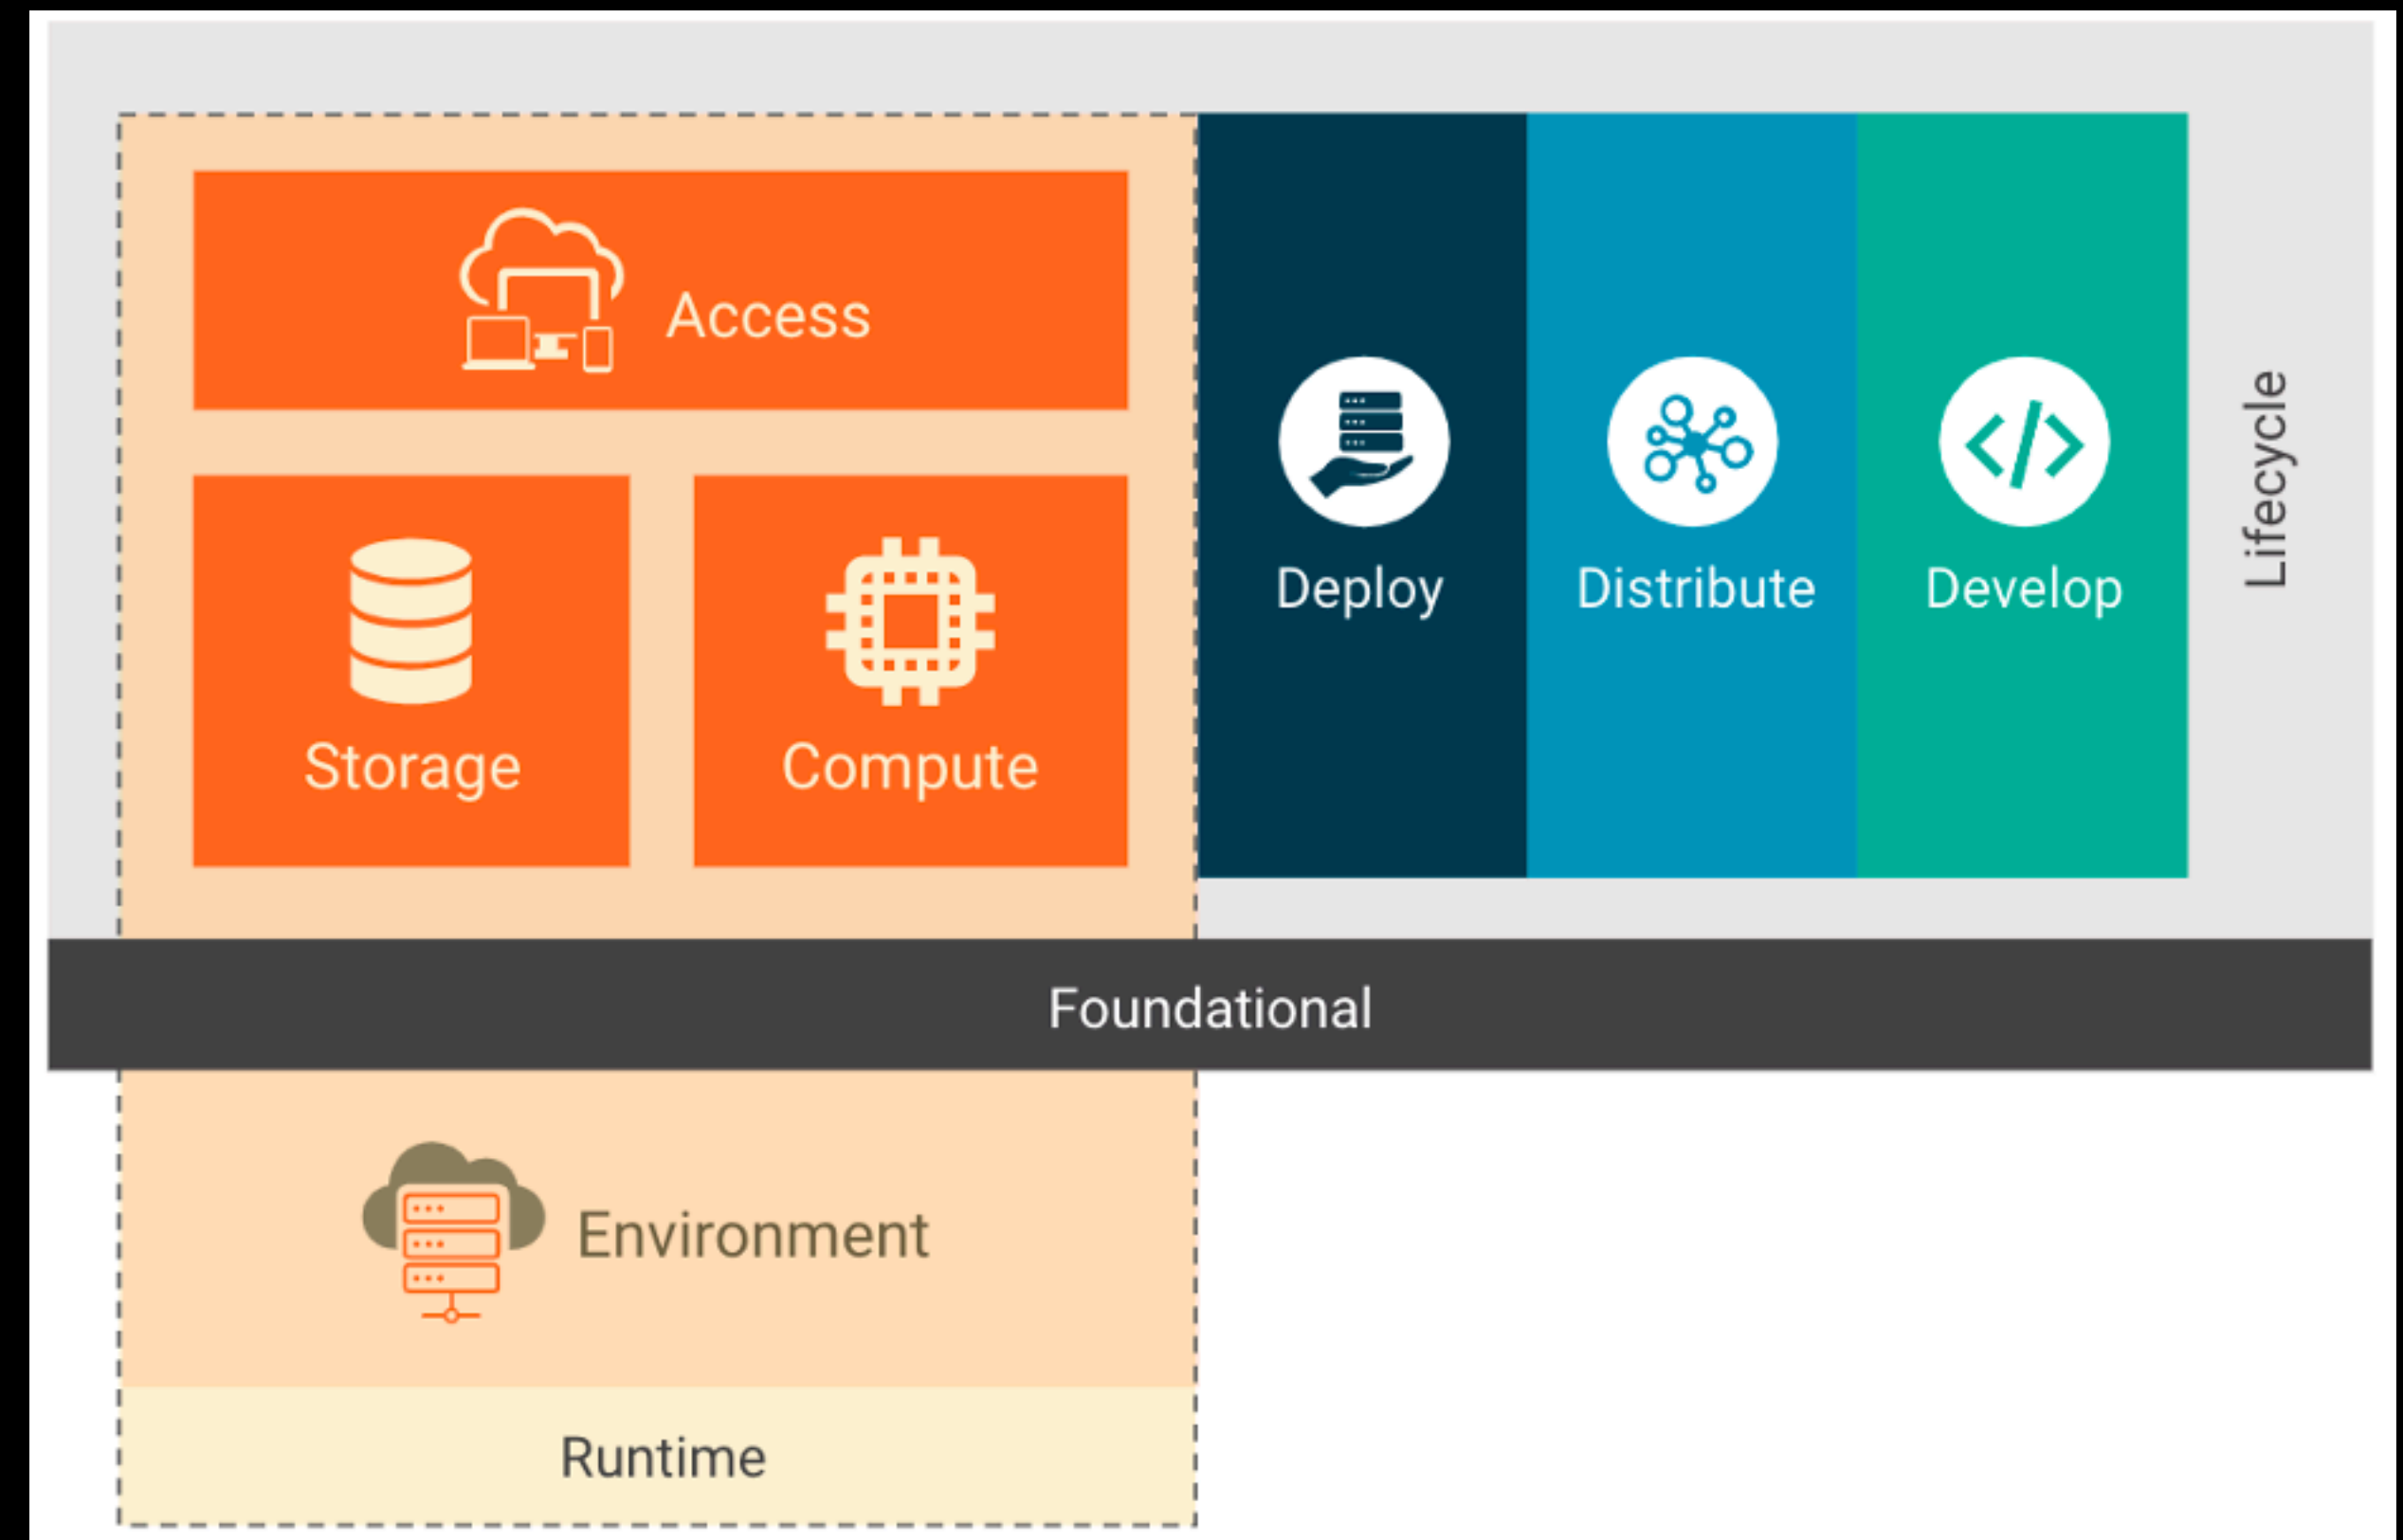


# How to do Kubernetes Security?

# Aspects of Kubernetes Security

## According to the Cloud Native Security Whitepaper

- Develop
- Distribute
- Deploy
- Runtime



# Implementing (Some) Security



Deploy

### Pre Flight Checks



Validate Image Integrity and Signature



Apply Image Runtime Policies



Apply Runtime Container Policies

### Runtime Policies



Apply Runtime Security Controls

- Standards based: NIST\*, CIS\*



Host Security

- Vulnerabilities
- Compliance Controls
- Micro-segmentation



Container Security

1. Workload Isolation
2. Network Policy
3. File Integrity
4. Process Integrity
5. Syscalls

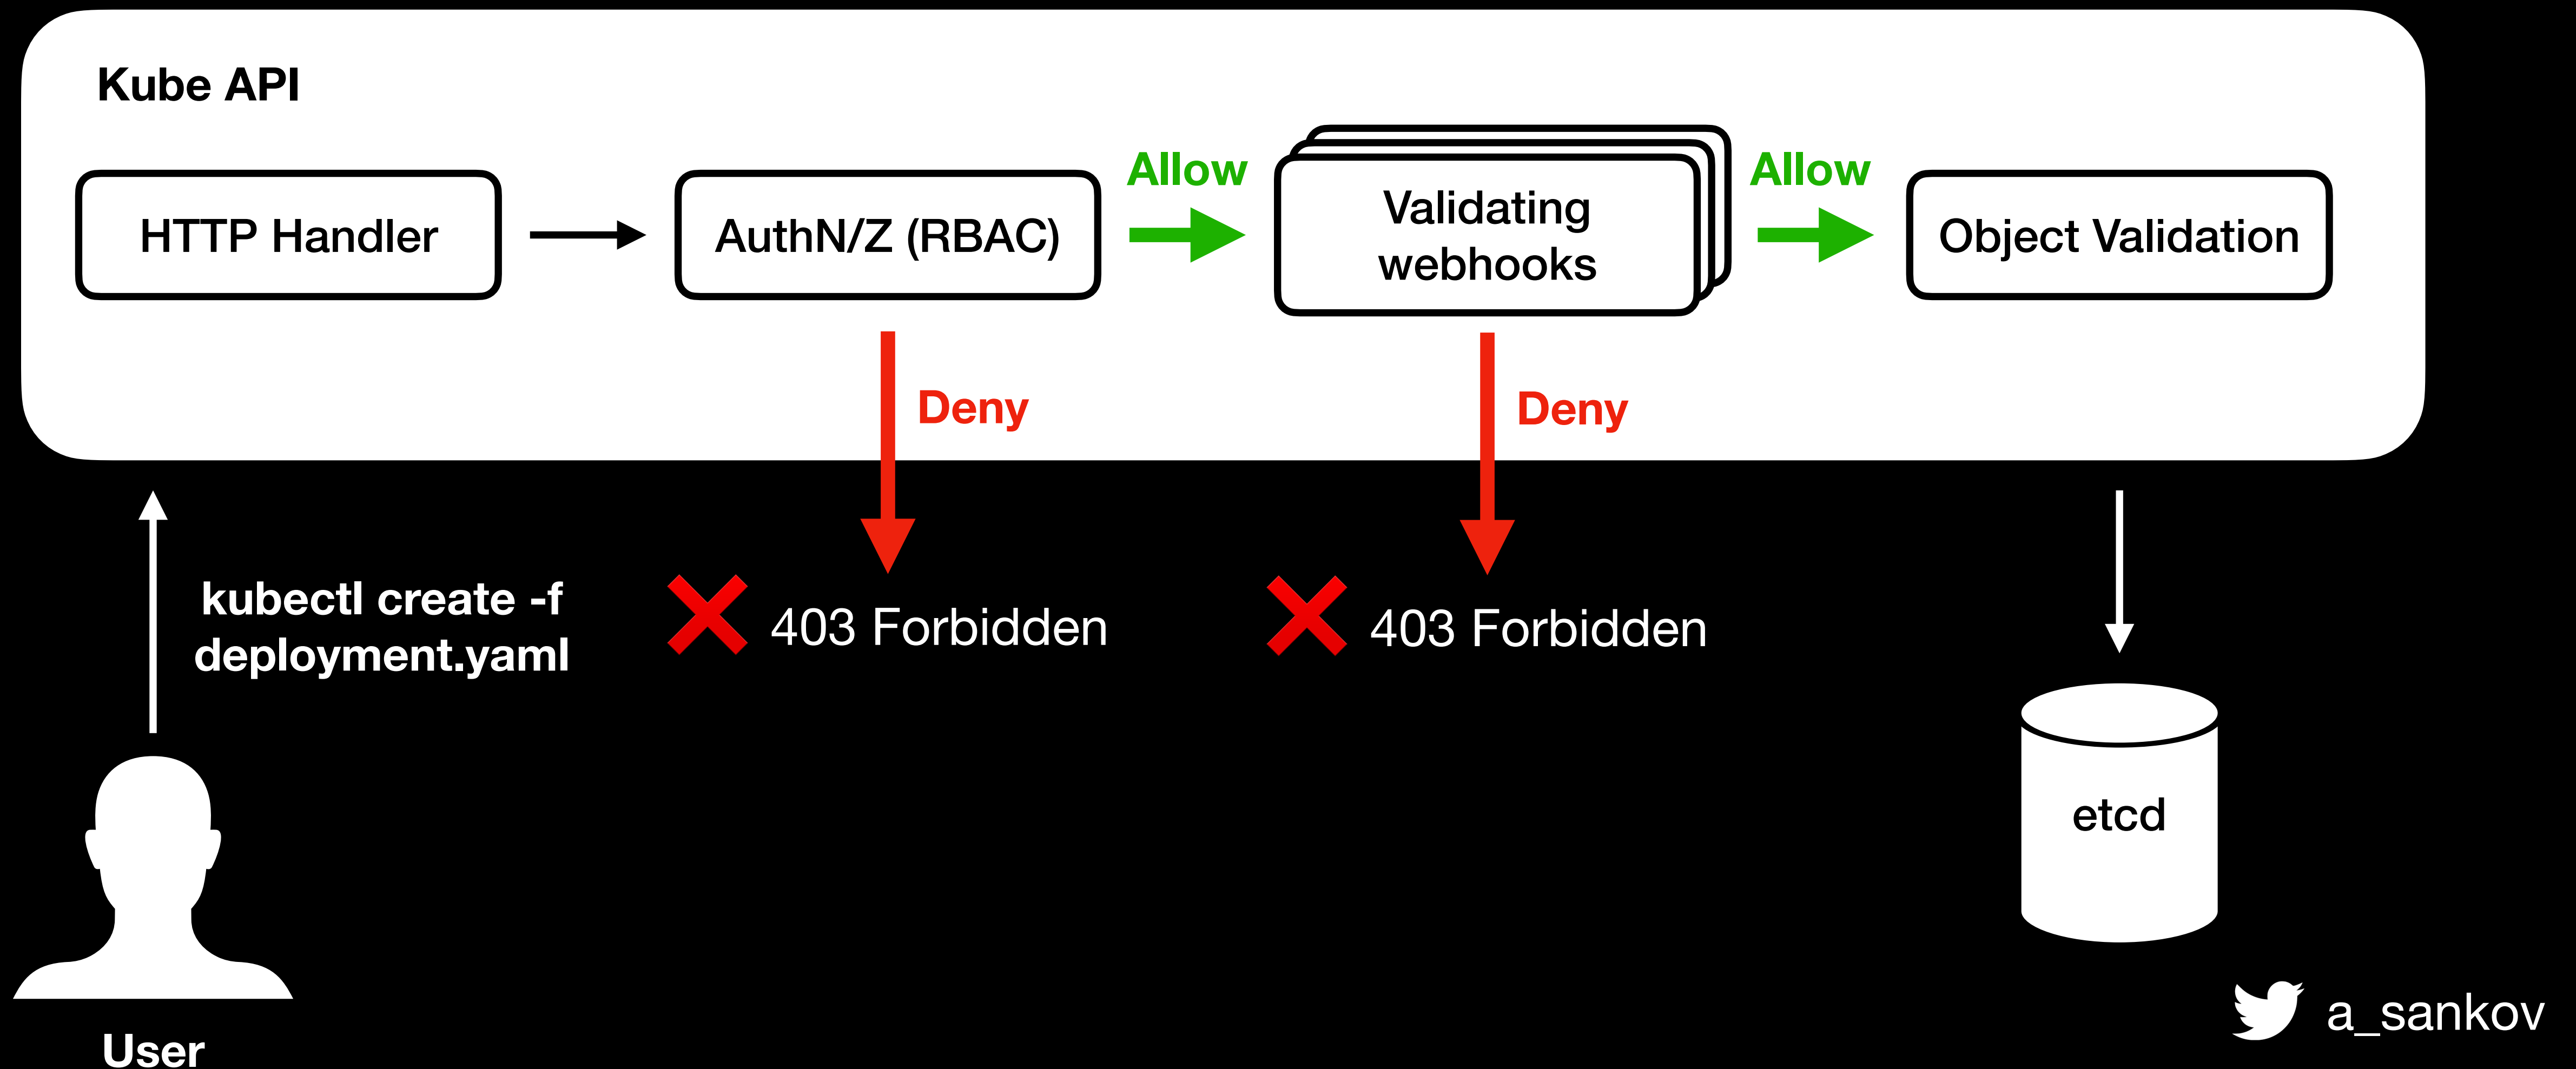
# Goals

- Enforce rules for our Kubernetes Deployments
- Rules are checked when Deployment object is applied to the API server
- Deployment object should not be created if it violates the rule

# Solution: Validating Webhooks

- Pluggable mechanism for adding additional verification to Kubernetes resource being created/updated
- Can have many of them, Kubernetes calls all in order
- If a validating webhook denies the request, Kubernetes aborts the operation
- Anyone can write and plug-in their own

# Kubernetes Validating Webhooks



# Solution: Validating Webhooks

```
apiVersion: admissionregistration.k8s.io/v1
kind: ValidatingWebhookConfiguration
metadata:
  name: "admission.asankov.dev"
webhooks:
- name: "admission.asankov.dev"
  rules:
  - apiGroups:    [""]
    apiVersions:  ["v1"]
    operations:   ["CREATE", "UPDATE"]
    resources:    ["Deployments"]
  clientConfig:
    url: "https://admission.asankov.dev/admission"
  admissionReviewVersions: ["v1", "v1beta1"]
  sideEffects: None
  timeoutSeconds: 5
```



# Solution: Validating Webhooks

```
apiVersion: admissionregistration.k8s.io/v1
kind: ValidatingWebhookConfiguration
metadata:
  name: "admission.asankov.dev"
webhooks:
- name: "admission.asankov.dev"
  rules:
  - apiGroups:  [""]
    apiVersions: ["v1"]
    operations:  ["CREATE", "UPDATE"]
    resources:   ["Deployments"]
  clientConfig:
    url: "https://admission.asankov.dev/admission"
  admissionReviewVersions: ["v1", "v1beta1"]
  sideEffects: None
  timeoutSeconds: 5
```

# Solution: Validating Webhooks

```
apiVersion: admissionregistration.k8s.io/v1
kind: ValidatingWebhookConfiguration
metadata:
  name: "admission.asankov.dev"
webhooks:
- name: "admission.asankov.dev"
  rules:
  - apiGroups:    [""]
    apiVersions:  ["v1"]
    operations:   ["CREATE", "UPDATE"]
    resources:    ["Deployments"]
  clientConfig:
    url: "https://admission.asankov.dev/admission"
  admissionReviewVersions: ["v1", "v1beta1"]
  sideEffects: None
  timeoutSeconds: 5
```

# Solution: Validating Webhooks

```
apiVersion: admissionregistration.k8s.io/v1
kind: ValidatingWebhookConfiguration
metadata:
  name: "admission.asankov.dev"
webhooks:
- name: "admission.asankov.dev"
  rules:
  - apiGroups:      [""]
    apiVersions:    ["v1"]
    operations:     ["CREATE", "UPDATE"]
    resources:      ["Deployments"]
  clientConfig:
    url: "https://admission.asankov.dev/admission"
  admissionReviewVersions: ["v1", "v1beta1"]
  sideEffects: None
  timeoutSeconds: 5
```


Kubernetes will call **this URL**  
when **Deployments**  
are being **created or updated**.

So... should I write my own  
Validating Webhook?

**Not necessarily.**

# Open Policy Agent (OPA)

- Open-source General-purpose policy agent
- Write policies in **Rego** language
- In: JSON input
- Out: JSON output
- Does not have anything to do with Kubernetes



# A (really) simple Rego policy

Input (JSON):

```
{  
  "conference": {  
    "name": "KubeHuddle"  
  }  
}
```

Output (JSON):

```
{  
  "allow": true  
}
```

Policy (Rego):

```
package policies  
  
default allow = false  
  
allow = true {  
    input.conference.name = "KubeHuddle"  
}
```

# A (really) simple Rego policy

Input (JSON):

```
{  
  "conference": {  
    "name": "SomeotherConf"  
  }  
}
```

Output (JSON):

```
{  
  "allow": false  
}
```

Policy (Rego):

```
package policies  
  
default allow = false  
  
allow = true {  
    input.conference.name = "KubeHuddle"  
}
```

# A (less) simple Rego policy

Input (JSON):

```
{
  "conference": {
    "name": "SomeotherConf",
    "venue": "SomeotherVenue"
  }
}
```

Output (JSON):

```
{
  "violation": [
    {"msg": "name and venue are wrong, - [SomeotherConf,
SomeotherVenue]"}
  ]
}
```

Policy (Rego):

```
package policies

violations[{"msg": msg}] {
  input.conference.name != "KubeHuddle"
  input.conference.venue != "Edinburgh"
  msg = sprintf("name and venue are wrong - [%s, %s]", [input.conference.name, input.conference.venue])
}
```



# Rego rules are just chained AND conditions

```
package policies

violations[{"msg": msg}] {
    input.conference.name != "KubeHuddle"
    input.conference.venue != "Edinburgh"
    msg = sprintf("name and venue are wrong - [%s, %s]", [input.conference.name, input.conference.venue])
}
```

## Translates to

```
if input.conference.name != "KubeHuddle" AND input.conference.venue != "Edinburgh" {
    msg = sprintf("name and venue are wrong - [%s, %s]", [input.conference.name, input.conference.venue])
    violations = append(violations, {"msg": msg})
}
```

**Which means that no message will be produced if conference.name is equal to “KubeHuddle” but the venue is different**

# A (less) simple Rego policy

Input (JSON):

```
{
  "conference": {
    "name": "SomeotherConf",
    "venue": "SomeotherVenue"
  }
}
```

Output (JSON):

```
{
  "violation": [
    {"msg": "name is wrong"},
    {"msg": "venue is wrong"}
  ]
}
```

Policy (Rego):

```
package policies

violations[{"msg": msg}] {
  input.conference.name != "KubeHuddle"
  msg := "name is wrong"
}

violations[{"msg": msg}] {
  input.conference.venue != "Edinburgh"
  msg := "venue is wrong"
}
```

# A (less) simple Rego policy

Input (JSON):

```
{
  "conference": {
    "name": "SomeotherConf",
    "venue": "Edinburgh"
  }
}
```

Output (JSON):

```
{
  "violation": [
    {"msg": "name is wrong"}
  ]
}
```

Policy (Rego):

```
package policies

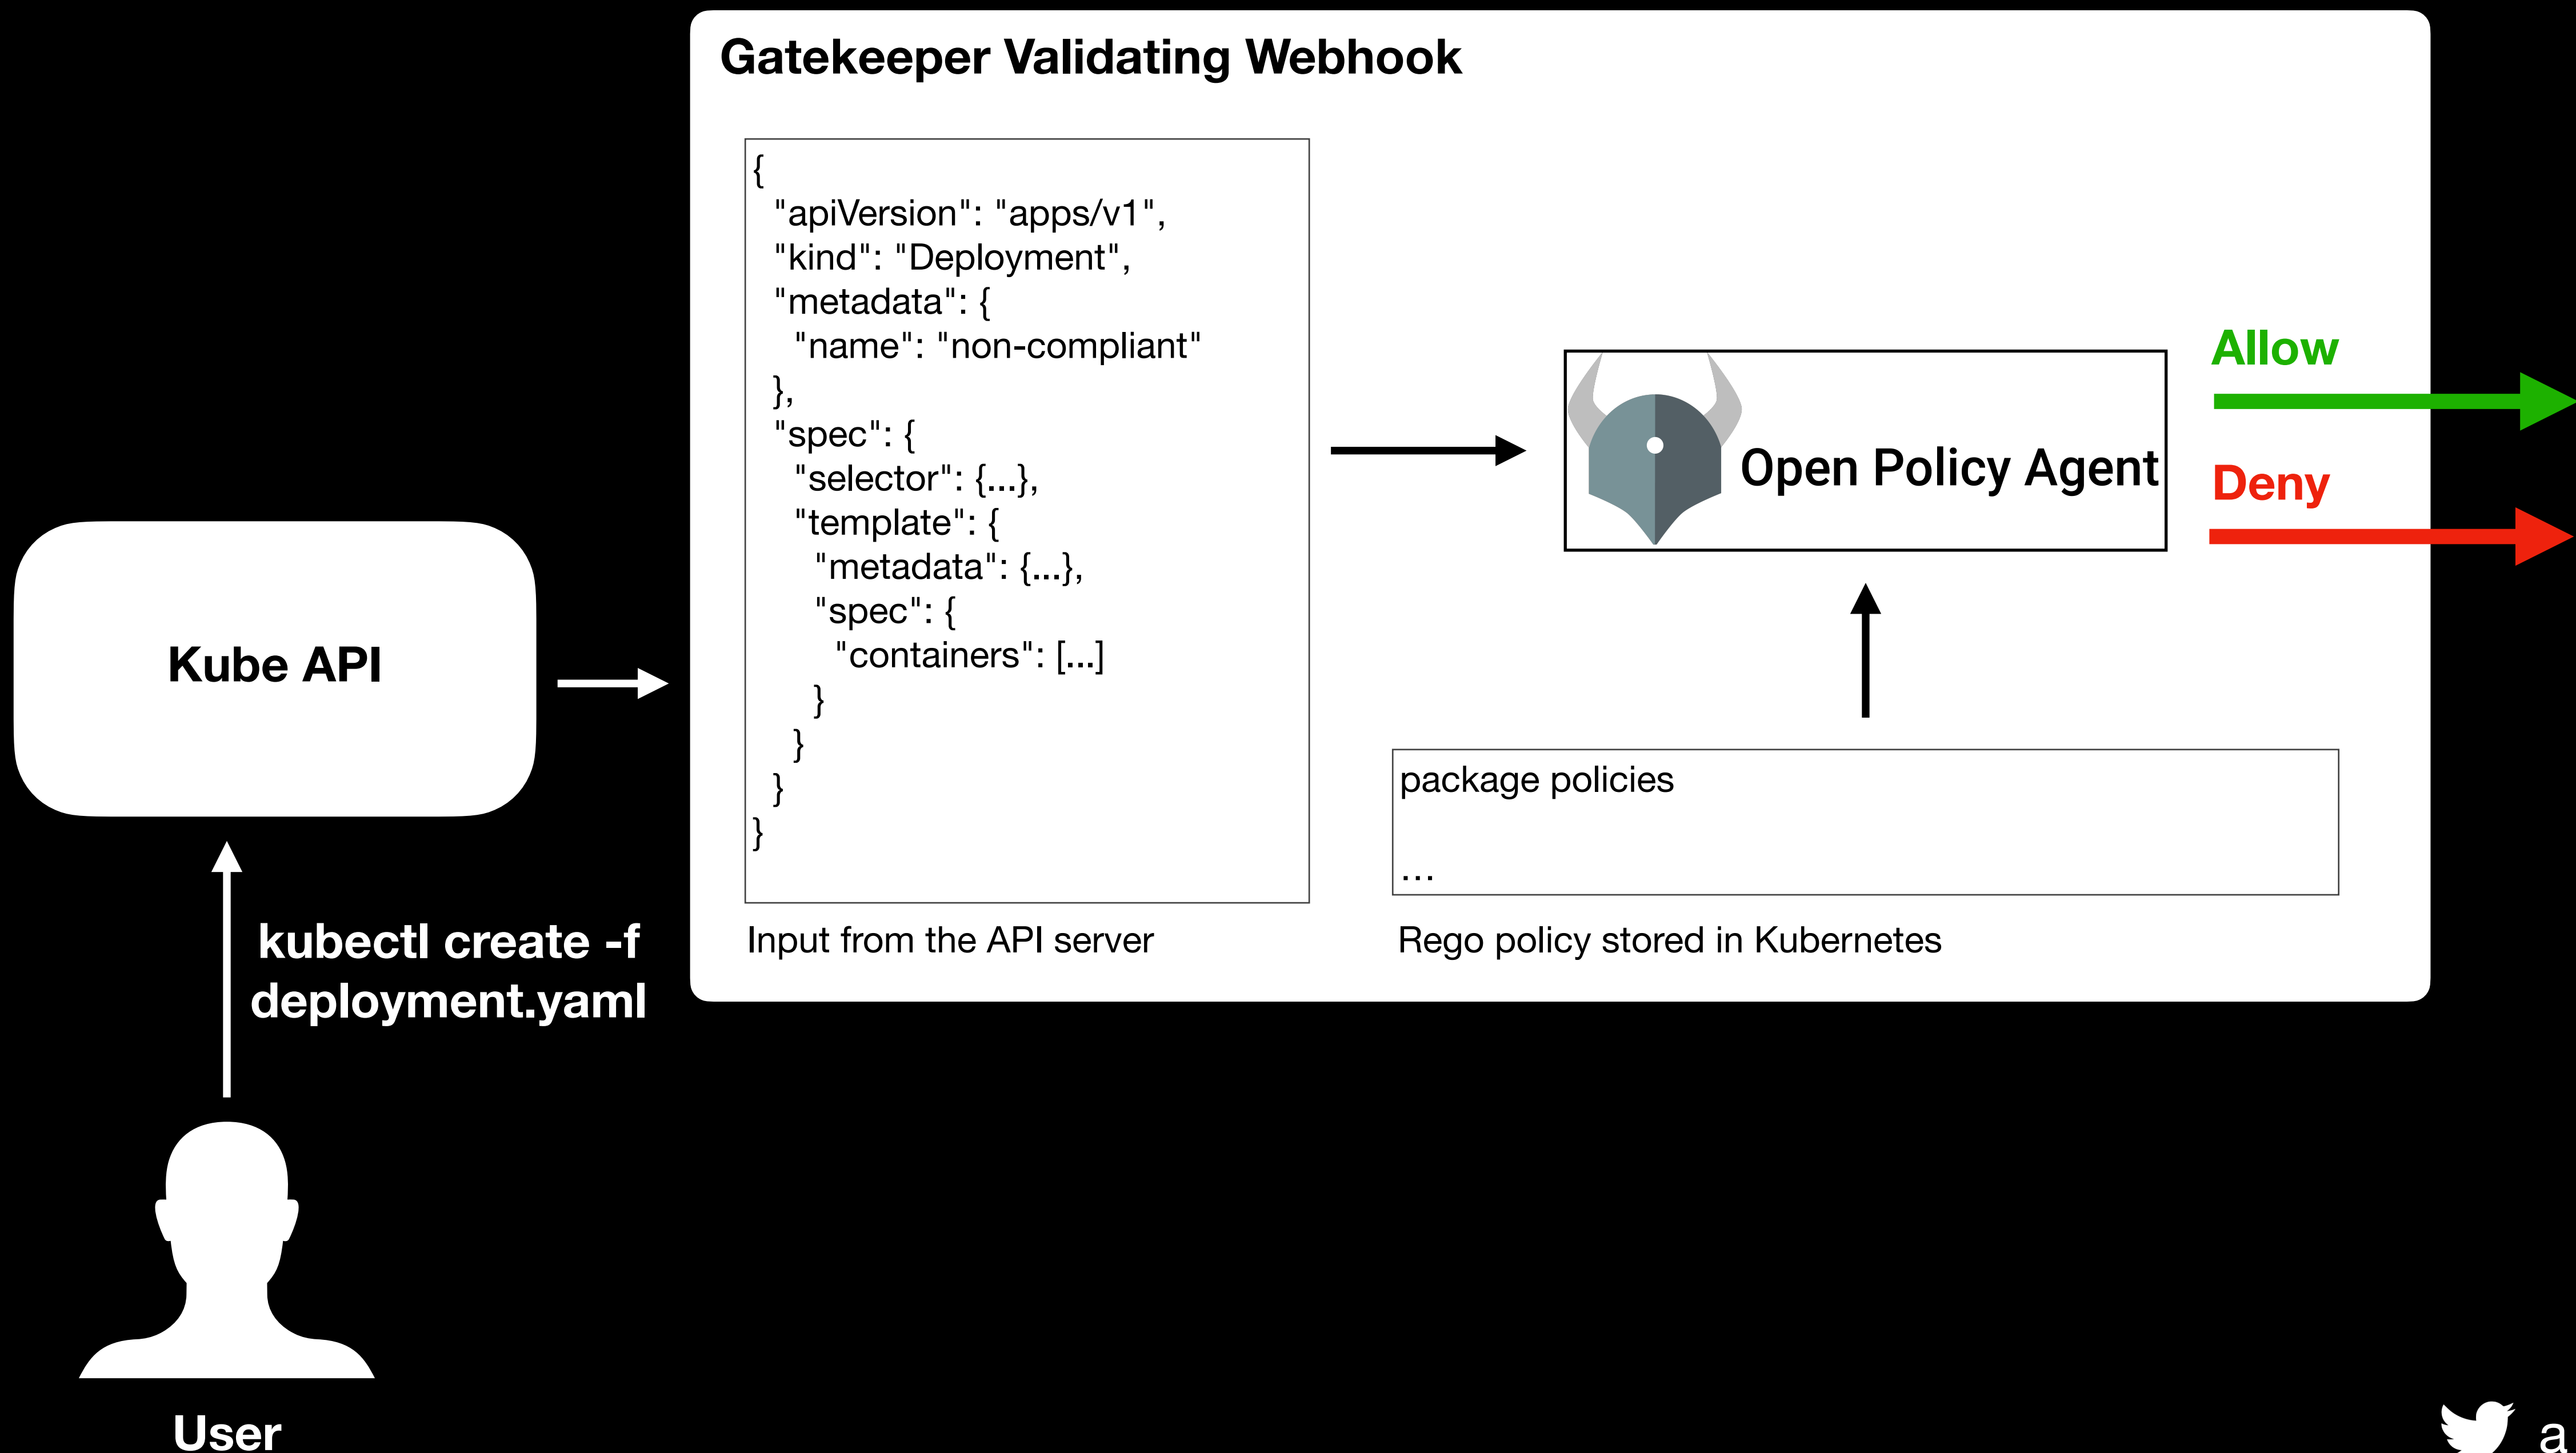
violations[{"msg": msg}] {
  input.conference.name != "KubeHuddle"
  msg := "name is wrong"
}

violations[{"msg": msg}] {
  input.conference.venue != "Edinburgh"
  msg := "venue is wrong"
}
```

**Let's add Kubernetes to the mix**

# OPA Gatekeeper

- First-class integration between OPA and Kubernetes
- Implements a validating webhook
- Calls OPA with the Kubernetes object as JSON input
- Returns a response that says whether the action can be completed based on the existing policies
- Policies are stored as Kubernetes objects (CRDs)



# Writing Gatekeeper policies

- Gatekeeper allows you to register policies as CRDs
- **ConstraintTemplate** - wraps the Rego policy
- **Constraint** - shows when to invoke the **ConstraintTemplate**

# Writing Gatekeeper policies

In programming terms

- **ConstraintTemplate** - a function that describes the policy
- **Constraint** - shows when to invoke the function



**Let's write some Gatekeeper  
Policies**

# Goals

## Reject Deployments that run as root

- A Deployment runs as Root if it has containers that:
  - Does not have a **securityContext**
  - OR does not have **securityContext.runAsUser** set
  - OR has **securityContext.runAsUser** set to **root uid (0)**

# Demo

# Alternatives

- Write your own Validation Webhook
- Kyverno
- Use PodSecurityPolicies/ PodSecurityStandards
- Use a proprietary solution

# Next steps

- Check out the links on the slides
- Other interesting talks about OPA:
  - <https://youtu.be/Vdy26oA3py8>
  - <https://youtu.be/ejH4EzmL7e0>
  - <https://youtu.be/RDWndems-sk>
- Write some policies

# Summary

- Kubernetes security is important
- Validating Webhooks are a pluggable mechanism for enforcing more granular rules on our Kubernetes objects
- OPA is a general-purpose policy agent
- Gatekeeper is Kubernetes-native OPA adapter
- Write rules and policies as code and interact with them the same you interact with other Kubernetes resources

Questions?

# Thank you!

 a\_sankov

 Anton Sankov

 [asankov.dev](https://asankov.dev)