

Securing Kubernetes with Open Policy Agent (OPA)

Anton Sankov, 19.06.2022

Anton Sankov


- Senior Software Engineer @ VMware Carbon Black
- Passionate about Kubernetes Security



a_sankov



Anton Sankov



<https://asankov.dev>



Agenda

Kubernetes

Kubernetes Security

Open Policy Agent (OPA) and Gatekeeper

Demo

**Every project mentioned in this
session is open-source**

Kubernetes



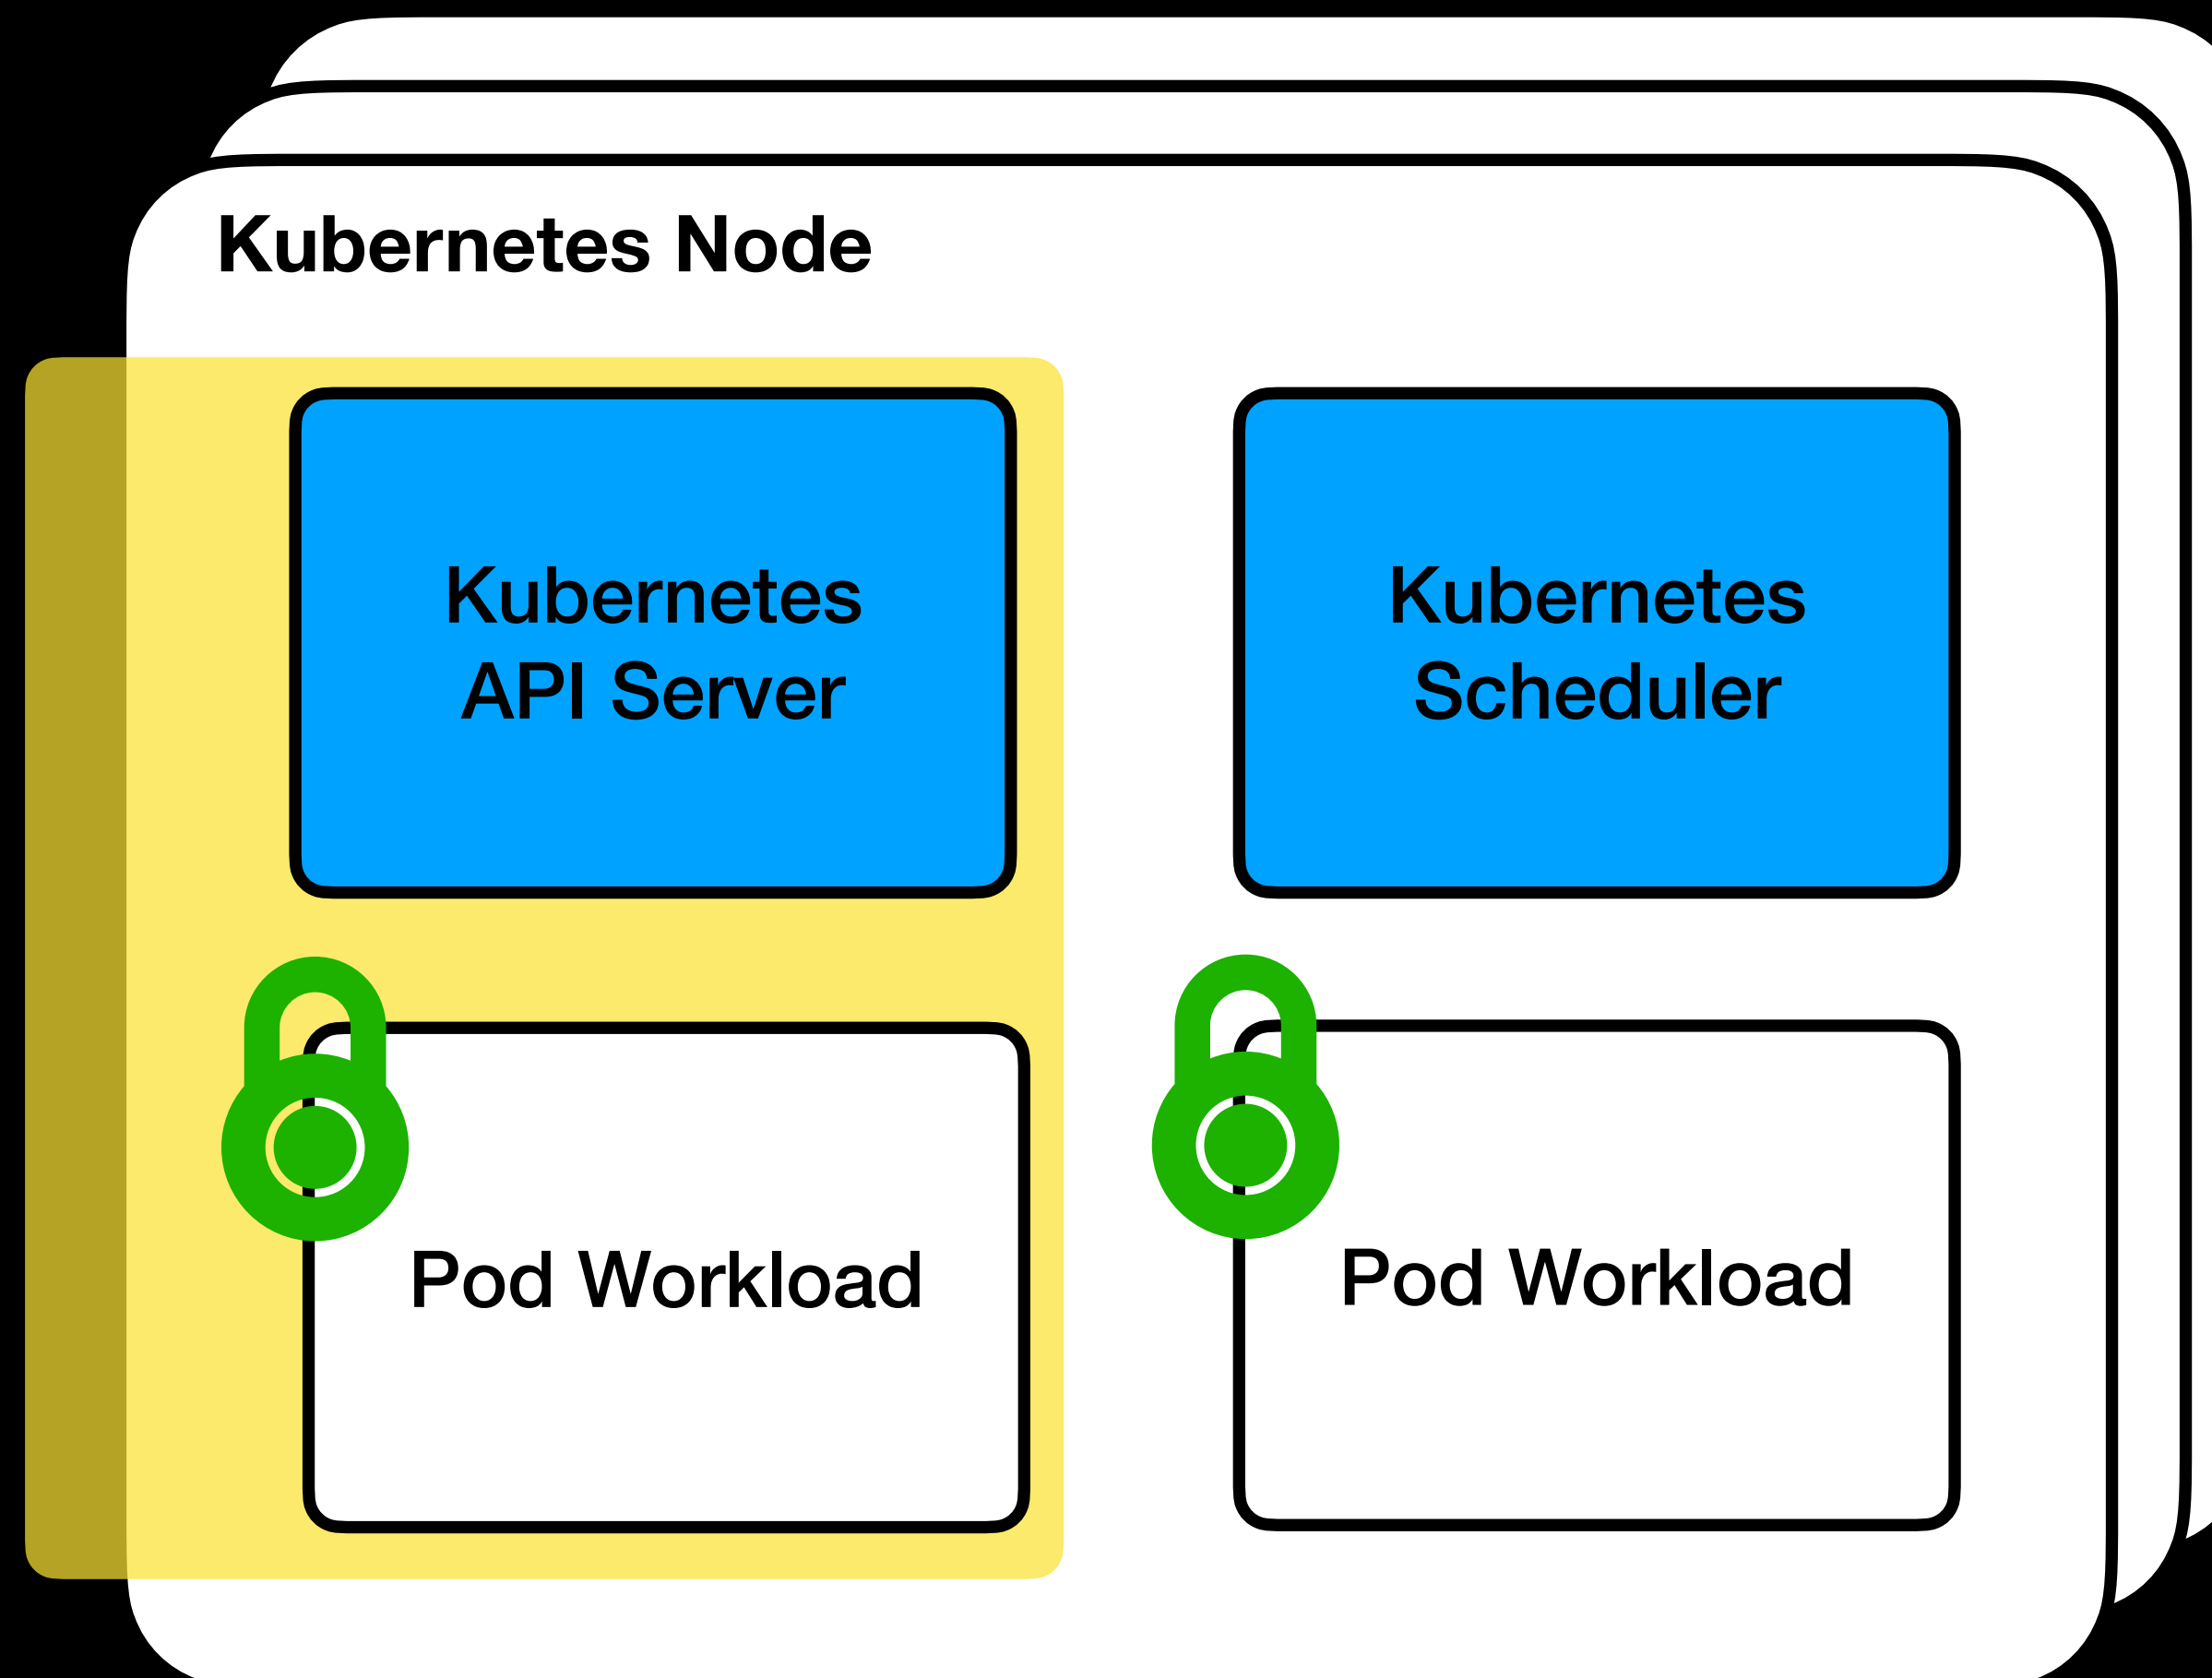
Open-source Container Orchestration tool

The de-facto standard for deploying applications in 2022

But this comes with some drawbacks

Bigger attack surface

**Publicly
accessible**



Kubernetes
Components

Custom
Workloads

Secured
Workloads

**But Kubernetes has build-in
security, right?**










Yes.

Kubernetes RBAC

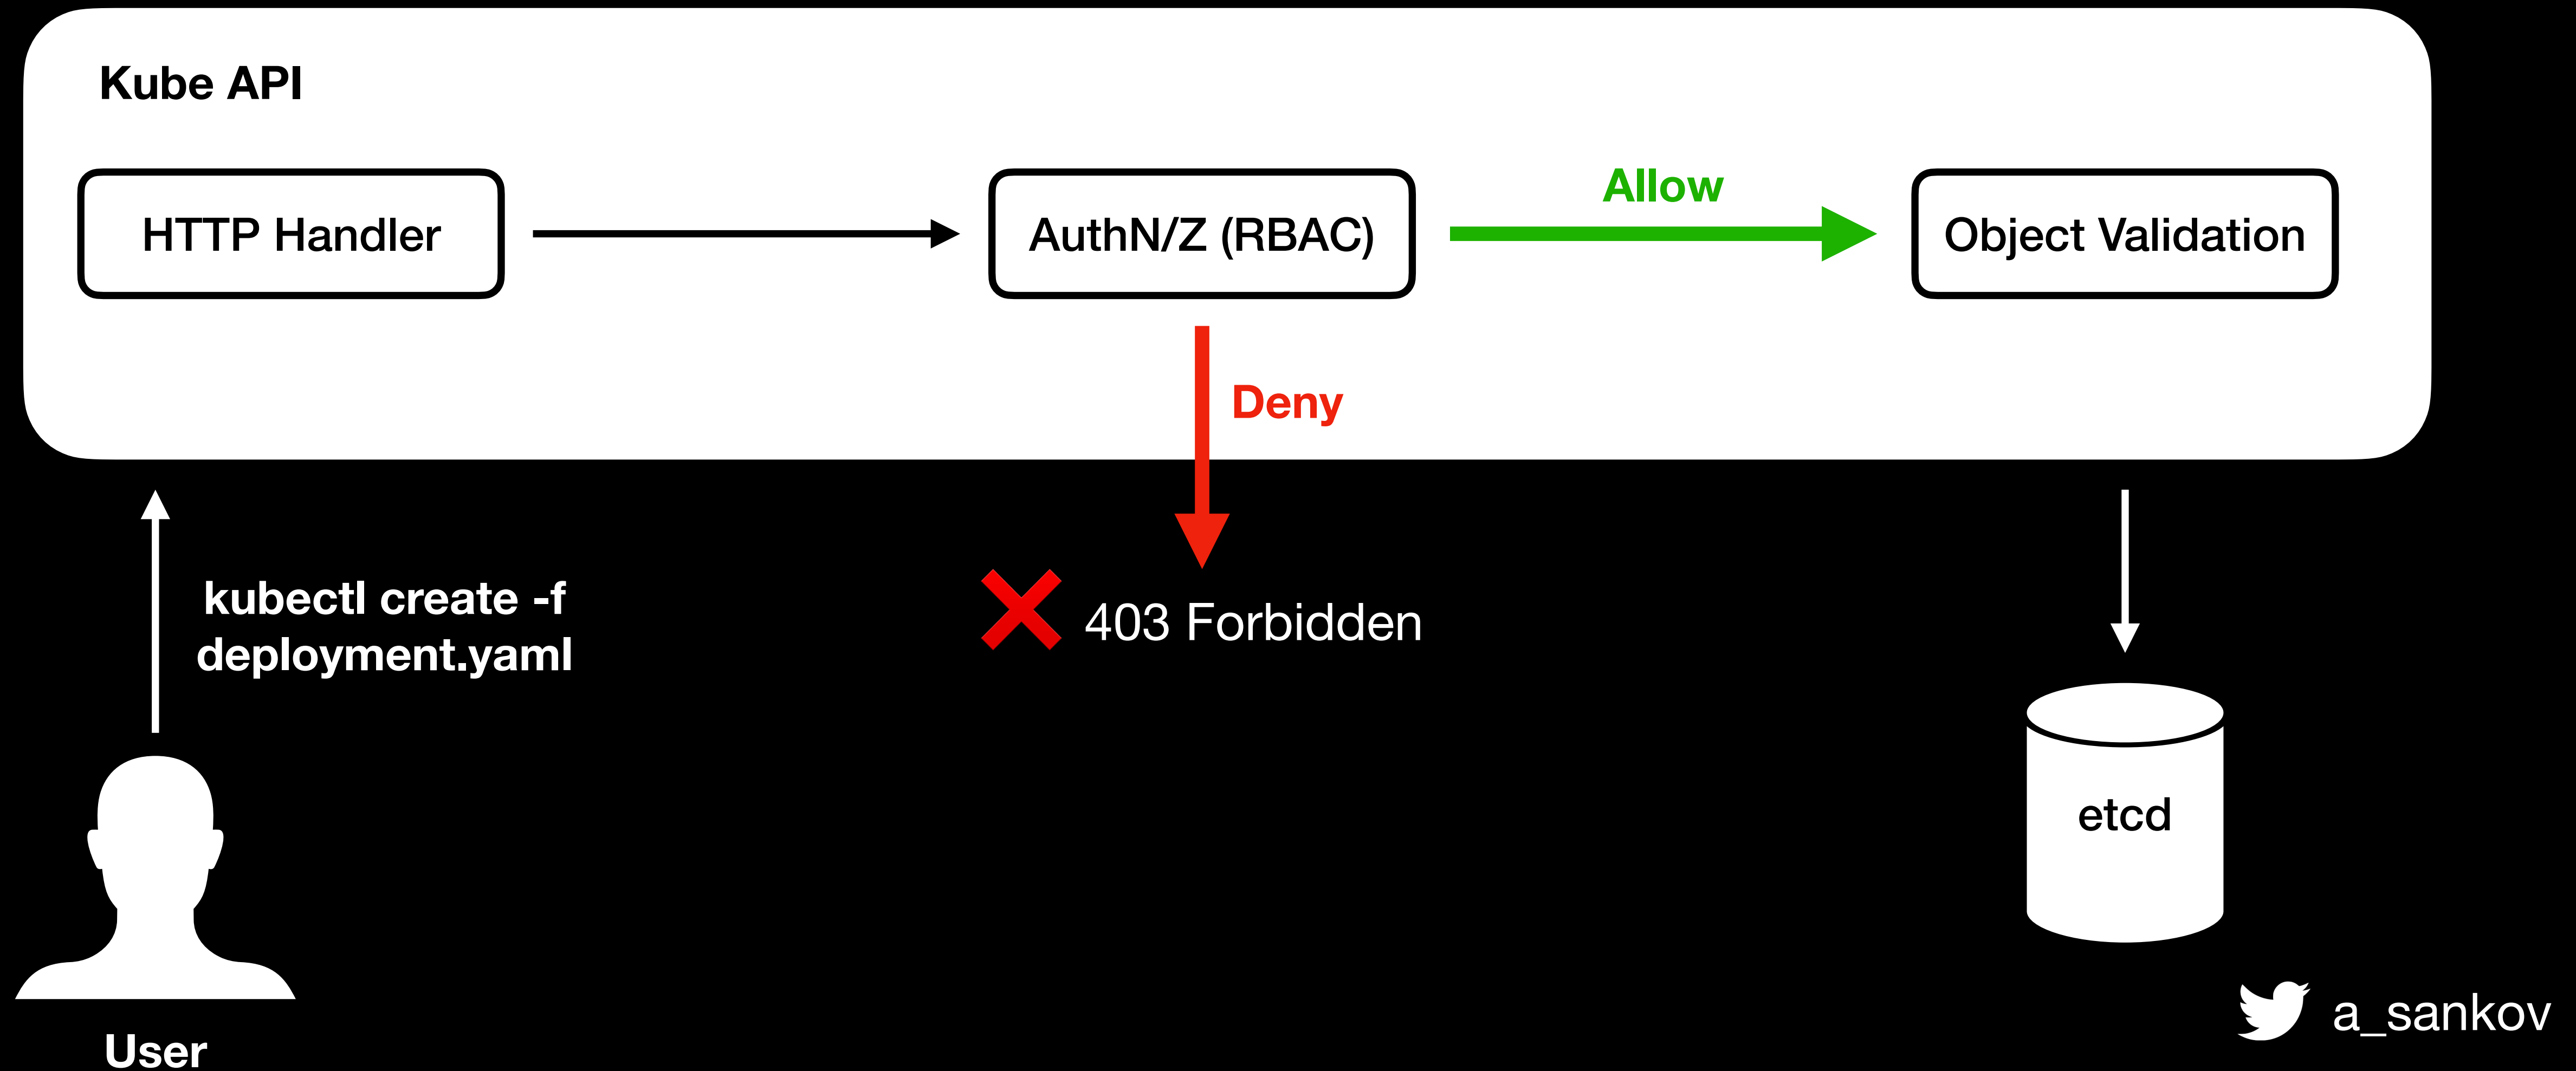
- Kubernetes has build-in RBAC (who can do what)
- Actions: **get, list, watch, delete, create, update**, etc.
- Resources: **Pods, deployments, services**, etc.

```
$ kubectl auth can-i create deployments  
yes
```


Kubernetes RBAC

	Get/List Deployments	Create/Update Deployments	Create RoleBindings
Developer			
DevOps / SRE			
Admin			

Kubernetes RBAC



Problem: No granularity

- A user that can create Deployments can create ANY kind of Deployments

Problem: No granularity

- Each organisation has rules that want to be enforced on the Kubernetes resources
- Examples:
 - All workloads should have team labels (for cost and ownership measuring)
 - All workloads should have resource limits (so that a rogue workload does not bring the whole cluster down)
 - Only images from trusted repositories should be used (so that an attacker cannot deploy a malicious image)

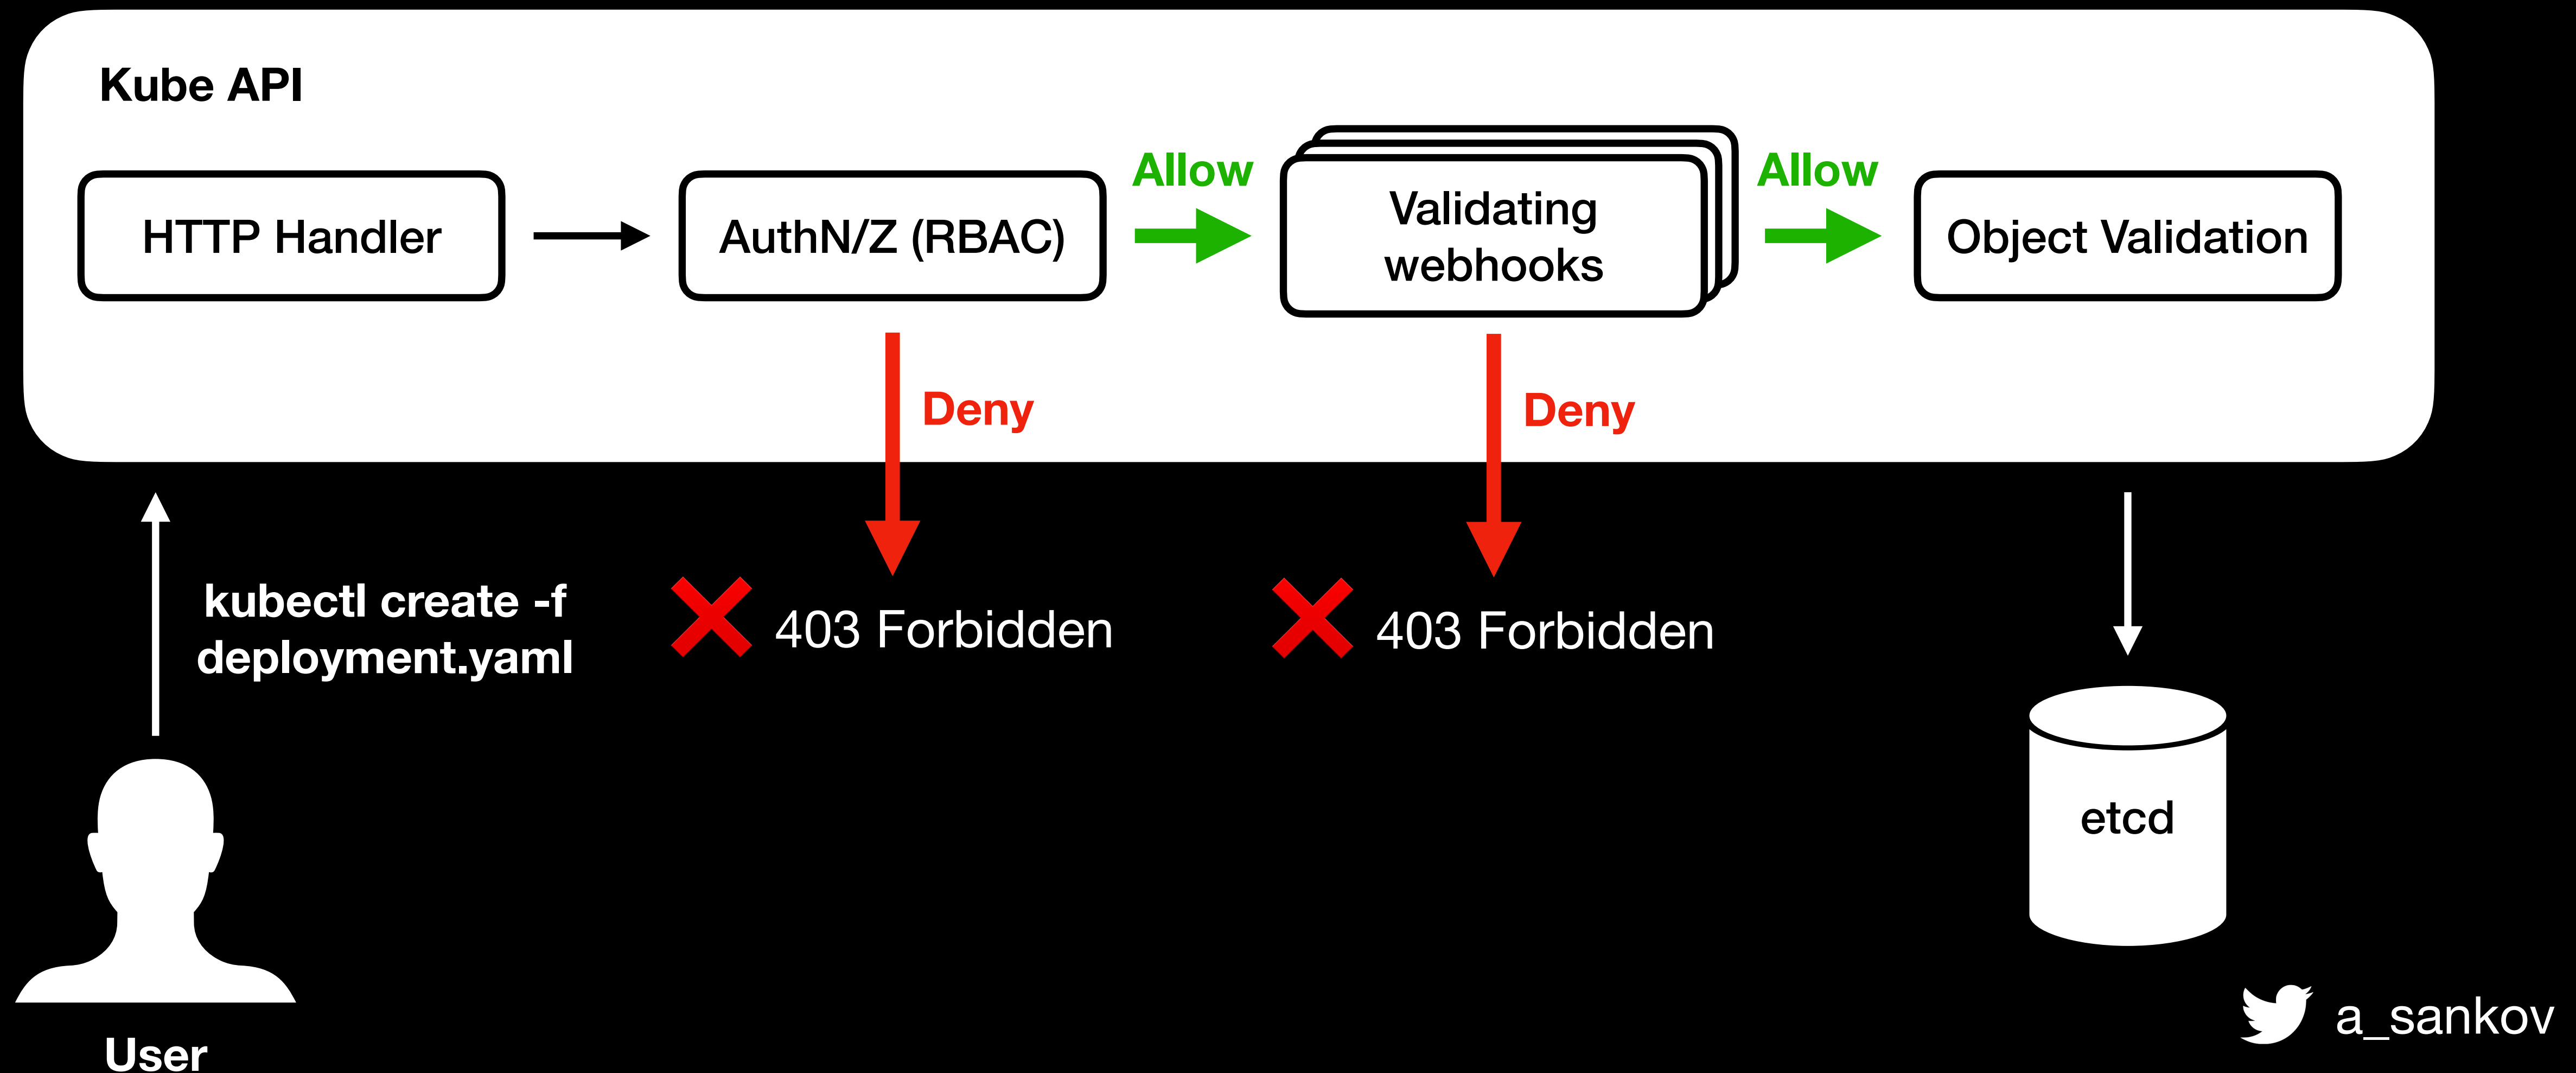
**But Kubernetes has build-in
solution for that problem, right?**

Sort of.

Solution: Validating Webhooks

- Pluggable mechanism for adding additional verification to Kubernetes resource being created/updated
- Can have many of them, Kubernetes calls all in order
- If a validating webhook denies the request, Kubernetes aborts the operation
- Anyone can write and plug-in their own

Kubernetes Validating Webhooks



Solution: Validating Webhooks

```
apiVersion: admissionregistration.k8s.io/v1
kind: ValidatingWebhookConfiguration
metadata:
  name: "admission.oscal.openlabs.cc"
webhooks:
- name: "admission.oscal.openlabs.cc"
  rules:
  - apiGroups:    [""]
    apiVersions:  ["v1"]
    operations:   ["CREATE"]
    resources:    ["Deployments"]
  clientConfig:
    url: "https://admission.oscal.openlabs.cc/admission"
  admissionReviewVersions: ["v1", "v1beta1"]
  sideEffects: None
  timeoutSeconds: 5
```


Solution: Validating Webhooks

```
apiVersion: admissionregistration.k8s.io/v1
kind: ValidatingWebhookConfiguration
metadata:
  name: "admission.oscal.openlabs.cc"
webhooks:
- name: "admission.oscal.openlabs.cc"
  rules:
  - apiGroups:    [""]
    apiVersions:  ["v1"]
    operations:   ["CREATE"]
    resources:    ["Deployments"]
  clientConfig:
    url: "https://admission.oscal.openlabs.cc/admission"
  admissionReviewVersions: ["v1", "v1beta1"]
  sideEffects: None
  timeoutSeconds: 5
```

Solution: Validating Webhooks

```
apiVersion: admissionregistration.k8s.io/v1
kind: ValidatingWebhookConfiguration
metadata:
  name: "admission.oscal.openlabs.cc"
webhooks:
- name: "admission.oscal.openlabs.cc"
  rules:
  - apiGroups:    [""]
    apiVersions:  ["v1"]
    operations:   ["CREATE"]
    resources:    ["Deployments"]
  clientConfig:
    url: "https://admission.oscal.openlabs.cc/admission"
  admissionReviewVersions: ["v1", "v1beta1"]
  sideEffects: None
  timeoutSeconds: 5
```

Solution: Validating Webhooks

```
apiVersion: admissionregistration.k8s.io/v1
kind: ValidatingWebhookConfiguration
metadata:
  name: "admission.oscal.openlabs.cc"
webhooks:
- name: "admission.oscal.openlabs.cc"
  rules:
  - apiGroups:      [""]
    apiVersions:    ["v1"]
    operations:     ["CREATE"]
    resources:      ["Deployments"]
  clientConfig:
    url: "https://admission.oscal.openlabs.cc/admission"
  admissionReviewVersions: ["v1", "v1beta1"]
  sideEffects: None
  timeoutSeconds: 5
```


Kubernetes will call **this URL**
when **Deployments** are being **created**.

So... should I write my own
Validating Webhook?

Not necessarily.

Open Policy Agent (OPA)

- Open-source General-purpose policy agent
- Write rules in Rego language
- In: JSON input
- Out: JSON output
- Does not have anything to do with Kubernetes



A (really) simple Rego policy

Input (JSON):

```
{  
  "conference": {  
    "name": "OSCAL"  
  }  
}
```

Output (JSON):

```
{  
  "allow": true  
}
```

Policy (Rego):

```
package bsidesdemo  
  
default allow = false  
  
allow = true {  
    input.conference.name = "OSCAL"  
}
```

A (really) simple Rego policy

Input (JSON):

```
{  
  "conference": {  
    "name": "SomeotherConf"  
  }  
}
```

Output (JSON):

```
{  
  "allow": false  
}
```

Policy (Rego):

```
package bsidesdemo  
  
default allow = false  
  
allow = true {  
    input.conference.name = "OSCAL"  
}
```

A (less) simple Rego policy

Input (JSON):

```
{
  "conference": {
    "name": "SomeotherConf",
    "venue": "SomeotherVenue"
  }
}
```

Output (JSON):

```
{
  "violation": [
    {"msg": "name and venue are wrong, - [SomeotherConf,
SomeotherVenue]"}
  ]
}
```

Policy (Rego):

```
package bsidesdemo

violations[{"msg": msg}] {
  input.conference.name != "OSCAL"
  input.conference.venue != "Tirana"
  msg = sprintf("name and venue are wrong - [%s, %s]", [input.conference.name, input.conference.venue])
}
```


Rego rules are just chained AND conditions

```
package bsidesdemo

violations[{"msg": msg}] {
    input.conference.name != "OSCAL"
    input.conference.venue != "Tirana"
    msg = sprintf("name and venue are wrong - [%s, %s]", [input.conference.name, input.conference.venue])
}
```

Translates to

```
if input.conference.name != "OSCAL" AND input.conference.venue != "Tirana" {
    msg = sprintf("name and venue are wrong - [%s, %s]", [input.conference.name, input.conference.venue])
    violations = append(violations, {"msg": msg})
}
```

Which means that no message will be produced if conference.name is equal to “OSCAL” but the venue is different

A (less) simple Rego policy

Input (JSON):

```
{
  "conference": {
    "name": "SomeotherConf",
    "venue": "SomeotherVenue"
  }
}
```

Output (JSON):

```
{
  "violation": [
    {"msg": "name is wrong"},
    {"msg": "venue is wrong"}
  ]
}
```

Policy (Rego):

```
package bsidesdemo

violations[{"msg": msg}] {
  input.conference.name != "OSCAL"
  msg := "name is wrong"
}

violations[{"msg": msg}] {
  input.conference.venue != "Tirana"
  msg := "venue is wrong"
}
```

A (less) simple Rego policy

Input (JSON):

```
{
  "conference": {
    "name": "SomeotherConf",
    "venue": "Tirana"
  }
}
```

Output (JSON):

```
{
  "violation": [
    {"msg": "name is wrong"}
  ]
}
```

Policy (Rego):

```
package bsidesdemo

violations[{"msg": msg}] {
  input.conference.name != "OSCAL"
  msg := "name is wrong"
}

violations[{"msg": msg}] {
  input.conference.venue != "Tirana"
  msg := "venue is wrong"
}
```

Good, but...

- Nothing so far was Kubernetes related
- The rules had some hard-coded values, which are not suitable for real environments

enter ... Gatekeeper

OPA Gatekeeper

- First-class integration between OPA and Kubernetes
- Implements a validating webhook
- Calls OPA with the Kubernetes object as JSON input
- Returns a response that says whether the action can be completed based on the existing policies
- Policies are stored as Kubernetes objects (CRDs)

Writing Gatekeeper policies

Gatekeeper represents Policies as Kubernetes objects
(CRDs)

ConstraintTemplate - describes the Rego rules
and the provided data

Constraint - shows how the **ConstraintTemplate**
should be enforced

Writing Gatekeeper policies

In programming terms:

ConstraintTemplate - a function that describes the policy, accepts arguments and returns a response

Constraint - shows how and when to invoke the function (what arguments to pass)

Writing Gatekeeper policies

```
apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate
metadata:
  name: k8srequiredlabels
spec:
  crd:
    spec:
      names:
        kind: K8sRequiredLabels
      validation:
        # Schema for the `parameters` field
        openAPIV3Schema:
          type: object
          properties:
            labels:
              type: array
              items:
                type: string
  targets:
    - target: admission.k8s.gatekeeper.sh
      rego: |
        package k8srequiredlabels

        violation[{"msg": msg, "details": {"missing_labels": missing}}] {
          provided := {label | input.review.object.metadata.labels[label]}
          required := {label | label := input.parameters.labels[_]}
          missing := required - provided
          count(missing) > 0
          msg := sprintf("you must provide labels: %v", [missing])
        }
```

```
apiVersion: constraints.gatekeeper.sh/v1beta1
kind: K8sRequiredLabels
metadata:
  name: deployments-must-have-gk
spec:
  match:
    kinds:
      - apiGroups: ["*"]
        kinds: ["Deployments"]
  parameters:
    labels: ["gatekeeper"]
```


Writing Gatekeeper policies

```
apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate
metadata:
  name: k8srequiredlabels
spec:
  crd:
    spec:
      names:
        kind: K8sRequiredLabels
      validation:
        # Schema for the `parameters` field
        openAPIV3Schema:
          type: object
          properties:
            labels:
              type: array
              items:
                type: string
  targets:
    - target: admission.k8s.gatekeeper.sh
      rego: |
        package k8srequiredlabels

        violation[{"msg": msg, "details": {"missing_labels": missing}}] {
          provided := {label | input.review.object.metadata.labels[label]}
          required := {label | label := input.parameters.labels[_]}
          missing := required - provided
          count(missing) > 0
          msg := sprintf("you must provide labels: %v", [missing])
        }
```

```
apiVersion: constraints.gatekeeper.sh/v1beta1
kind: K8sRequiredLabels
metadata:
  name: deployments-must-have-gk
spec:
  match:
    kinds:
      - apiGroups: ["*"]
        kinds: ["Deployments"]
  parameters:
    labels: ["gatekeeper"]
```

**Tells Kubernetes to invoke this rule
ONLY when Deployments are being created.
It will not be invoked for any other kind.**

Writing Gatekeeper policies

```
apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate
metadata:
  name: k8srequiredlabels
spec:
  crd:
    spec:
      names:
        kind: K8sRequiredLabels
      validation:
        # Schema for the `parameters` field
        openAPIV3Schema:
          type: object
          properties:
            labels:
              type: array
              items:
                type: string
  targets:
    - target: admission.k8s.gatekeeper.sh
      rego: |
        package k8srequiredlabels
```

```
violation[{"msg": msg, "details": {"missing_labels": missing}}] {
  provided := {label | input.review.object.metadata.labels[label]}
  required := {label | label := input.parameters.labels[_]}
  missing := required - provided
  count(missing) > 0
  msg := sprintf("you must provide labels: %v", [missing])
}
```

```
apiVersion: constraints.gatekeeper.sh/v1beta1
kind: K8sRequiredLabels
metadata:
  name: deployments-must-have-gk
spec:
  match:
    kinds:
      - apiGroups: ["*"]
        kinds: ["Deployments"]
  parameters:
    labels: ["gatekeeper"]
```

Parametrizes some parameters in the rule so that we can reuse ConstraintTemplates by create new Constraints (much like we reuse function)

Writing Gatekeeper policies

```
apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate
metadata:
  name: k8srequiredlabels
spec:
  crd:
    spec:
      names:
        kind: K8sRequiredLabels
      validation:
        # Schema for the `parameters` field
        openAPIV3Schema:
          type: object
          properties:
            labels:
              type: array
              items:
                type: string
  targets:
    - target: admission.k8s.gatekeeper.sh
      rego: |
        package k8srequiredlabels
```

```
violation[{"msg": msg, "details": {"missing_labels": missing}}] {
  provided := {label | input.review.object.metadata.labels[label]}
  required := {label | label := input.parameters.labels[_]}
  missing := required - provided
  count(missing) > 0
  msg := sprintf("you must provide labels: %v", [missing])
}
```

```
apiVersion: constraints.gatekeeper.sh/v1beta1
kind: K8sRequiredLabels
metadata:
  name: deployments-must-have-gk
spec:
  match:
    kinds:
      - apiGroups: ["*"]
        kinds: ["Deployments"]
  parameters:
    labels: ["gatekeeper"]
```

**The Kubernetes object being created.
Input provided by Gatekeeper**

Demo

Alternatives

- Just use RBAC
- Lower the visibility of your Kubernetes as much as possible and hope someone does not get into your private network
- Use PodSecurityPolicies/ PodSecurityStandards
- Use a proprietary solution

Next steps

- Check out the links on the slides
- Other interesting talks about OPA:
 - <https://youtu.be/Vdy26oA3py8>
 - <https://youtu.be/ejH4EzmL7e0>
 - <https://youtu.be/RDWndems-sk>
- Go write some policies

<https://asankov.dev/k8s-sec-opa/>

Summary

- Build-in Kubernetes security (RBAC) is not enough for most organisations
- Validating Webhooks are a pluggable mechanism for enforcing more granular rules on our Kubernetes objects
- OPA is a general-purpose policy agent
- Gatekeeper is Kubernetes-native OPA adapter
- Write rules and policies as code and interact with them the same you interact with other Kubernetes resources

Questions?

Thank you!

 a_sankov

 Anton Sankov

 asankov.dev